

# Machine Learning in Robotics

## Lecture 4: Unsupervised Clustering

**Prof. Dongheui Lee**

*Institute of Automatic Control Engineering  
Technische Universität München*

dhlee@tum.de

# Today Lecture Outline

- Unsupervised Clustering
- Similarity measures
- Criterion functions
- Iterative optimization algorithms
- K-means & Variations
- Hierarchical clustering

# Supervised vs. Unsupervised learning

- Supervised learning
  - A pattern is a pair of variables  $\{x, \omega\}$  where  $x$  is a collection of observations or features (feature vector) and  $\omega$  is the concept behind the observation (label)
- Unsupervised learning
  - Use unlabeled data, a collection of feature vectors without the class label  $\omega$

# Supervised vs. Unsupervised learning


- Supervised learning
  - A pattern is a pair of variables  $\{x, \omega\}$  where  $x$  is a collection of observations or features (feature vector) and  $\omega$  is the concept behind the observation (label)
- Unsupervised learning
  - Use unlabeled data, a collection of feature vectors without the class label  $\omega$
  - These methods are called unsupervised because they are not provided the correct answer

# Supervised vs. Unsupervised learning


- Supervised learning
  - A pattern is a pair of variables  $\{x, \omega\}$  where  $x$  is a collection of observations or features (feature vector) and  $\omega$  is the concept behind the observation (label)
- Unsupervised learning
  - Use unlabeled data, a collection of feature vectors without the class label  $\omega$
  - These methods are called unsupervised because they are not provided the correct answer
  - Unsupervised methods may appear to have limited capabilities, but they are useful
    - ▶ Labeling large data sets can be a costly procedure
    - ▶ Class labels may not be known beforehand
    - ▶ Large datasets can be compressed by finding a small set of proto-types

# TEST : Unsupervised clustering

What is the natural grouping among those objects?



**Many possibilities!!!**



# Two approaches for unsupervised learning

- Parametric approaches
  - Functional forms for the underlying class-conditional densities are assumed, and we must estimate the parameters

$$p(\mathbf{x}|\theta) = \sum_{i=1}^K p(\mathbf{x}|\omega_i, \theta_i) p(\omega_i)$$

# Two approaches for unsupervised learning

- Parametric approaches
  - Functional forms for the underlying class-conditional densities are assumed, and we must estimate the parameters

$$p(\mathbf{x}|\theta) = \sum_{i=1}^K p(\mathbf{x}|\omega_i, \theta_i) p(\omega_i)$$

- Non-parametric approaches
  - No assumptions are made about the underlying densities
  - Instead, seek a partition of the data into clusters
  - These methods are typically referred to as clustering



# Vector Quantization (VQ)

- Vector Quantization is a lossy data compression method
- Mapping  $n$  feature vectors  $X = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$  to  $K$  classes of feature vectors  $Y = \{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(K)}\}$

$$\mathbf{y}^{(j)} = c(\mathbf{x}^{(i)}), \quad j = 1, \dots, K; \quad i = 1, \dots, n; \quad K < n$$

- Code vector  $\mathbf{y}^{(j)}$ 
  - Cluster centers
  - Points with high density in feature space
- Code book  $Y = \{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(K)}\}$ 
  - The set of all code vectors



# Nonparametric clustering

Nonparametric clustering involves three steps:

- Defining a measure of (dis)similarity between examples
- Defining a criterion (Distortion) function for clustering
- Defining an algorithm to minimize (or maximize) a criterion (distortion) function

# Similarity Measures

- A measuring rule of  $d(x, y)$  for the distance between two vectors  $x$  and  $y$  is considered a metric if it satisfies the following properties.
  - non-negativity:  $d(x, y) \geq 0$
  - reflexivity:  $d(x, y) = 0$  if and only if  $x = y$
  - symmetry:  $d(x, y) = d(y, x)$
  - triangle inequality:  $d(x, y) + d(y, z) \geq d(x, z)$

# Similarity Measures

- The most general form of distance metric is the power norm

$$d(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^m |x_i - y_i|^q \right)^{1/q}$$

- where  $q \geq 1$  is a selectable parameter - general Minkowski metric
- When  $q = 2$ , Euclidean metric  $d(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\| = \sqrt{(\mathbf{x} - \mathbf{y})^T (\mathbf{x} - \mathbf{y})}$
- When  $q = 1$ , Manhattan or city block metric - sum of the absolute distances along each of the  $m$  coordinate axes.
- Notice that the above distance metrics are measures of dissimilarity.

# Similarity Measures

- Inner product

$$s(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

- When the features are binary-valued (0 or 1), the normalized inner product is a measure of the relative possession of common attributes.
- One variation is Tanimoto distance, a ratio of the number of shared attributes to the number possessed by  $\mathbf{x}$  or  $\mathbf{y}$ .

$$s(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}^T \mathbf{y}}{\mathbf{x}^T \mathbf{x} + \mathbf{y}^T \mathbf{y} - \mathbf{x}^T \mathbf{y}}$$

# Similarity measure



# Criterion (Distortion) function

- Once a (dis)similarity measure has been determined, we need to define a criterion function to be optimized.
- This function measures the clustering quality of any partition of the data.
  - The most widely used criterion function for clustering is the *sum-of-square-error*

$$J = \sum_{i=1}^K \sum_{x \in \omega_i} (x - y^{(i)})^2 \quad \text{where } y^{(i)} = \frac{1}{N_i} \sum_{x \in \omega_i} x$$

Here,  $N_i$  is the number of samples in the  $\omega_i$

- This criterion measures how well the data set is represented by the cluster centers
- Other criterion functions exist, based on the scatter matrices used in Linear Discriminant Analysis (LDA).
  - Trace criterion  $tr[S_w]$
  - Determinant criterion  $|S_w|$
  - Invariant criterion  $tr[S_T^{-1} S_w]$

# Iterative optimization

- Once a criterion function has been defined, we must find a partition of the data set that minimizes the criterion.
- Exhaustive enumeration of all partitions, which guarantees the optimal solution, is unfeasible.
- Common approach is to proceed in an iterative fashion
  - Find reasonable initial partition
  - Move samples from one cluster to another in order to reduce the criterion function
- These iterative methods produce sub-optimal solution but are computationally tractable



# Iterative optimization

- Once a criterion function has been defined, we must find a partition of the data set that minimizes the criterion.
- Exhaustive enumeration of all partitions, which guarantees the optimal solution, is unfeasible.
- Common approach is to proceed in an iterative fashion
  - Find reasonable initial partition
  - Move samples from one cluster to another in order to reduce the criterion function
- These iterative methods produce sub-optimal solution but are computationally tractable
- Approaches
  - K-means algorithm
  - LBG algorithm
  - Fuzzy K-means algorithm
  - Basic iterative minimum-squared-error clustering

# k-means, EM algorithm

- EM algorithm
  - Useful when estimating an optimal solution of a problem including hidden information
  - An iterative method which alternates between performing an expectation (E) step and a maximization (M) step
  - E step: computes the expectation of the log-likelihood evaluated using the current estimate for the latent variables
  - M step: computes parameters maximizing the expected log-likelihood found on the E step
  - Local convergence, No guarantee for global convergence
- K-means
  - A simple example of the EM optimization algorithm

# k-means algorithm (Lloyd, 1982)

Cluster dataset  $X = \{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$ , given  $K$

1. Initialization: Choose  $K$  random vectors  $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$  as an initial mean set

# k-means algorithm (Lloyd, 1982)

Cluster dataset  $X = \{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$ , given  $K$

1. Initialization: Choose  $K$  random vectors  $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$  as an initial mean set
2. E-step: For each data point, find the closest class and label them. Dataset is divided into  $K$  classes

$$X_i = \{x^{(j)} \mid d(x^{(j)}, y^{(i)}) \leq d(x^{(j)}, y^{(k)}), j = 1, \dots, n; k = 1, \dots, K\}$$



# k-means algorithm (Lloyd, 1982)

Cluster dataset  $X = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$ , given  $K$

1. Initialization: Choose  $K$  random vectors  $\mathbf{Y} = \{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(K)}\}$  as an initial mean set
2. E-step: For each data point, find the closest class and label them. Dataset is divided into  $K$  classes

$$X_i = \{\mathbf{x}^{(j)} \mid d(\mathbf{x}^{(j)}, \mathbf{y}^{(i)}) \leq d(\mathbf{x}^{(j)}, \mathbf{y}^{(k)}), j = 1, \dots, n; k = 1, \dots, K\}$$

3. M-step: From the current clusters, their mean vectors are updated


$$\mathbf{y}^{(i)} = \frac{1}{N_i} \sum_{\mathbf{x} \in \omega_i} \mathbf{x} = \mathcal{C}(X_i)$$

4. Calculate the total distortion, the sum of the distance between each datapoint and its closest cluster mean.


$$J = \sum_{i=1}^K \sum_{\mathbf{x} \in \omega_i} (\mathbf{x} - \mathbf{y}^{(i)})^2$$

5. Evaluate the convergence. If converged, stop. Else, go to step 2.

# Illustration of k-means algorithm



K=2



Blue : after E-step

Red : after M-step

Source: C. M. Bishop, Pattern Recognition and Machine Learning, Springer, 2006

# k-means clustering applications



Source: C. M. Bishop, Pattern Recognition and Machine Learning, Springer, 2006



# k-means clustering: Remarks

- The way to initialize the means is not specified. One popular way to start is to randomly choose  $k$  of the samples.
- A local optimization algorithm
  - Converge to a local rather than global minimum of  $J$
  - Convergence properties of the k-means algorithm [MacQueen 1967]
  - The results depend on the initial values for the means. The standard solution is to try a number of different starting points.
- Uniform search
- Slow convergence
- It can happen that the set of samples closest to  $\mathbf{y}^{(i)}$  is empty, so that  $\mathbf{y}^{(i)}$  cannot be updated.
- The results depend on the value of  $K$ .

# Basic Iterative Minimum-Squared-Error Clustering

- A sequential version of the k-means clustering algorithm.
  - k-means procedure waits until all  $n$  samples have been reclassified before updating.
  - the Basic Iterative Minimum-Squared-Error Clustering updates after each sample is reclassified.
- Disadvantages
  - more susceptible to being trapped in local minima
  - results depend on the order in which the candidates are selected
- Merits
  - at least a stepwise optimal procedure
  - suitable to problems in which samples are acquired sequentially and clustering must be done online

# Fuzzy k-means algorithm

- Allows one piece of data to belong to two or more clusters
- Minimizing the below objective function

$$J = \sum_{k=1}^K \sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m (\mathbf{x}^{(i)} - \mathbf{y}^{(k)})^2$$

- $m$  is a parameter to adjust the blending of different clusters. If  $m = 0$ , k-means. For  $m \geq 1$ , the criterion allows belonging to multiple clusters
- $P(\omega_k | \mathbf{x}^{(i)})$  is the degree of membership of  $\mathbf{x}^{(i)}$  in the cluster  $k$
- an iterative optimization of the objective function shown above, with the update of membership  $P(\omega_k | \mathbf{x}^{(i)})$  and the cluster centers  $\mathbf{y}^{(k)}$  by:

$$P(\omega_k | \mathbf{x}^{(i)}) = \frac{1}{\sum_{j=1}^K \left( \frac{\|\mathbf{x}^{(i)} - \mathbf{y}^{(k)}\|}{\|\mathbf{x}^{(i)} - \mathbf{y}^{(j)}\|} \right)^{\frac{2}{m-1}}}$$
$$\mathbf{y}^{(k)} = \frac{\sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m \mathbf{x}^{(i)}}{\sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m}$$

# Fuzzy k-means algorithm : Pseudocode

Cluster dataset, given  $K$

1. Initialization: Initialize  $K$  random mean vectors  $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$ . Initialize  $P(\omega_k | x^i)$  for  $i = 1..n, k = 1..K$
2. Normalize  $P(\omega_k | x^{(i)})$  so that  $\sum_k^K P(\omega_k | x^{(i)}) = 1$  for  $i = 1..n$

# Fuzzy k-means algorithm : Pseudocode

Cluster dataset, given  $K$

1. Initialization: Initialize  $K$  random mean vectors  $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$ . Initialize  $P(\omega_k | x^i)$  for  $i = 1..n, k = 1..K$
2. Normalize  $P(\omega_k | x^{(i)})$  so that  $\sum_k^K P(\omega_k | x^{(i)}) = 1$  for  $i = 1..n$
3. Update mean vectors

$$y^{(k)} = \frac{\sum_{i=1}^n P(\omega_k | x^{(i)})^m x^{(i)}}{\sum_{i=1}^n P(\omega_k | x^{(i)})^m}$$

4. Update the degree of membership of  $x^{(i)}$  in the cluster  $k$

$$P(\omega_k | x^{(i)}) = \frac{1}{\sum_{j=1}^K \left( \frac{\|x^{(i)} - y^{(k)}\|}{\|x^{(i)} - y^{(j)}\|} \right)^{\frac{2}{m-1}}}$$

# Fuzzy k-means algorithm : Pseudocode

Cluster dataset, given  $K$

1. Initialization: Initialize  $K$  random mean vectors  $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$ . Initialize  $P(\omega_k | x^i)$  for  $i = 1..n, k = 1..K$
2. Normalize  $P(\omega_k | x^{(i)})$  so that  $\sum_k^K P(\omega_k | x^{(i)}) = 1$  for  $i = 1..n$
3. Update mean vectors

$$y^{(k)} = \frac{\sum_{i=1}^n P(\omega_k | x^{(i)})^m x^{(i)}}{\sum_{i=1}^n P(\omega_k | x^{(i)})^m}$$

4. Update the degree of membership of  $x^{(i)}$  in the cluster  $k$

$$P(\omega_k | x^{(i)}) = \frac{1}{\sum_{j=1}^K \left( \frac{\|x^{(i)} - y^{(k)}\|}{\|x^{(i)} - y^{(j)}\|} \right)^{\frac{2}{m-1}}}$$


5. If  $\max_{i,k} P(\omega_k | x^{(i)})$  is converged, then stop. Else, go to step 2.

## Another variation of k-means

- ISODATA, which stands for Iterative Self-Organizing Data Analysis Technique (Algorithm) is an extension to the k-means algorithm with some heuristics to automatically select the number of clusters
- The algorithm works in an iterative fashion
  - (1) Perform k-means clustering
  - (2) Split any clusters whose samples are sufficiently dissimilar
  - (3) Merge any two clusters sufficiently close
  - (4) Go to (1)

# Hierarchical clustering


- k-means and ISODATA create disjoint clusters, resulting in a "flat" data representation
  - ▶ However, sometimes it is desirable to obtain a hierarchical representation of data, with clusters and sub-clusters arranged in a tree-structured fashion (i.e., biological taxonomy)
- Hierarchical clustering methods can be grouped in two general classes
  - ▶ Agglomerative (bottom-up, merging) : Starting with  $n$  singleton clusters, successively merge clusters until one cluster is left
  - ▶ Divisive (top-down, splitting) : Starting with a unique cluster, successively split the clusters until  $n$  singleton examples are left





# Dendrograms

- The preferred representation for hierarchical clusters is the dendrogram
- The dendrogram is a binary tree that shows the structure of the clusters
  - ▶ In addition to the binary tree, the dendrogram provides the similarity measure between clusters (the vertical axis)
- An alternative representation is based on sets
  - ▶  $\{\{x_1, \{x_2, x_3\}\}, \{\{\{x_4, x_5\}, \{x_6, x_7\}\}, x_8\}\}$
  - ▶ However, unlike the dendrogram, sets cannot express quantitative information



# Divisive clustering

1. Start with one large cluster
  2. Find "worst" cluster
  3. Split it
  4. If  $K < n$ , go to step 2
- How to choose the "worst" cluster
  - How to split clusters
  - The computations required by divisive clustering are more intensive than for agglomerative clustering methods.

# Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

# Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Choose a class which has the largest distortion among current classes

$$J_i \geq J_j, \quad \forall j = 1, \dots, k$$

# Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Choose a class which has the largest distortion among current classes

$$J_i \geq J_j, \quad \forall j = 1, \dots, k$$

3. Split the class into two subclasses by using a small random vector

$$X_a = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} + \mathbf{v}_i \right\}$$

$$X_b = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} - \mathbf{v}_i \right\}$$

# Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Choose a class which has the largest distortion among current classes

$$J_i \geq J_j, \quad \forall j = 1, \dots, k$$

3. Split the class into two subclasses by using a small random vector

$$\mathbf{X}_a = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} + \mathbf{v}_i \right\}$$

$$\mathbf{X}_b = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} - \mathbf{v}_i \right\}$$

4. Update the code vectors

$$\mathbf{y}^{(i)} := \mathcal{C}(\mathbf{X}_a) \quad \text{and} \quad \mathbf{y}^{(k+1)} := \mathcal{C}(\mathbf{X}_b)$$

# Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Choose a class which has the largest distortion among current classes

$$J_i \geq J_j, \quad \forall j = 1, \dots, k$$

3. Split the class into two subclasses by using a small random vector

$$\mathbf{X}_a = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} + \mathbf{v}_i \right\}$$


$$\mathbf{X}_b = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} - \mathbf{v}_i \right\}$$

4. Update the code vectors


$$\mathbf{y}^{(i)} := \mathcal{C}(\mathbf{X}_a) \quad \text{and} \quad \mathbf{y}^{(k+1)} := \mathcal{C}(\mathbf{X}_b)$$

5. If  $k = K$ , stop. Else,  $k := k + 1$  and go to step 2.

# k-means vs. Non-uniform binary split



(a) k-means



(b) Binary split

Non-uniform binary split: Not good quality of codebook. But very fast



# LBG algorithm

- Combination of k-means and Binary split
- Proposed by Linde, Buzo, and Gray (1980)
- Instead of random initial codebook, let's use codebook from binary split
- Faster convergence and better quality codebook than k-means algorithm

# Pseudo-code of LBG algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

# Pseudo-code of LBG algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Splitting: For each cluster, split it into two subclasses.  $k := 2k$

$$\mathbf{y}_a^{(i)} = \mathbf{y}^{(i)} + \mathbf{v}_i \text{ and } \mathbf{y}_b^{(i)} = \mathbf{y}^{(i)} - \mathbf{v}_i$$

# Pseudo-code of LBG algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Splitting: For each cluster, split it into two subclasses.  $k := 2k$

$$\mathbf{y}_a^{(i)} = \mathbf{y}^{(i)} + \mathbf{v}_i \text{ and } \mathbf{y}_b^{(i)} = \mathbf{y}^{(i)} - \mathbf{v}_i$$

3. Iteration (k-means algorithm)

- a E-step: For each data point, find the closest cluster which achieves the minimum distance measure
- b M-step: From the current cluster, update their mean vectors
$$\mathbf{y}^{(i)} = \mathcal{C}(\mathbf{X}_i)$$
- c Iterate E-step and M-step until it converges

# Pseudo-code of LBG algorithm

1. Initialization: Calculate a center for all data points.  $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Splitting: For each cluster, split it into two subclasses.  $k := 2k$


$$\mathbf{y}_a^{(i)} = \mathbf{y}^{(i)} + \mathbf{v}_i \text{ and } \mathbf{y}_b^{(i)} = \mathbf{y}^{(i)} - \mathbf{v}_i$$

3. Iteration (k-means algorithm)


- a E-step: For each data point, find the closest cluster which achieves the minimum distance measure
- b M-step: From the current cluster, update their mean vectors  
 $\mathbf{y}^{(i)} = \mathcal{C}(\mathbf{X}_i)$
- c Iterate E-step and M-step until it converges

4. If the desired number of code vectors is obtained ( $k = K$ ), stop. Else, go to step 2.


# Comparison



**kmeans**



**Binary Split**



**LBG**

# Agglomerative clustering

1. Start with  $n$  singleton cluster
  2. Find nearest clusters
  3. Merge them
  4. If  $K > 1$ , go to step 2
- How to find the "nearest" pair of clusters
    - ▶ Minimum distance  $d_{min}(X_i, X_j) = \min_{\mathbf{x} \in \omega_i, \mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
    - ▶ Maximum distance  $d_{max}(X_i, X_j) = \max_{\mathbf{x} \in \omega_i, \mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
    - ▶ Average distance  $d_{avg}(X_i, X_j) = \frac{1}{N_i N_j} \sum_{\mathbf{x} \in \omega_i} \sum_{\mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
    - ▶ Mean distance  $d_{mean}(X_i, X_j) = \|\mathbf{y}^{(i)} - \mathbf{y}^{(j)}\|$

# Agglomerative clustering

## Minimum distance

- When  $d_{min}$  is used to measure distance between clusters, the algorithm is called the *nearest neighbor* or *single-linkage* clustering algorithm
- If the algorithm is allowed to run until only one cluster remains, the result is a minimum spanning tree (MST)

## Maximum distance

- When  $d_{max}$  is used to measure distance between clusters, the algorithm is called the *farthest neighbor* or *complete-linkage* clustering algorithm
- From a graph-theoretic point of view, each cluster constitutes a complete sub-graph

## Average and mean distance


- The average and mean distance approaches are more robust to outliers
- Of the two, the mean distance is computationally more attractive, since Notice that the average distance approach involves the computation of  $N_i N_j$  distances for each pair of clusters



# Agglomerative clustering example

Perform agglomerative clustering on the following dataset using the single-linkage metric

- $X = \{1, 3, 4, 9, 10, 13, 21, 23, 28, 29\}$
- In case of ties, always merge the pair of clusters with the largest mean
- Indicate the order in which the merging operations occur





# Multi-Robot Exploration

- Task definition
  - ▶ Exploration of unknown areas by means of  $k$  mobile robots.
  - ▶ Application: Search/Rescue robot, planetary exploration, reconnaissance
- Problems
  - ▶ Reduce the difference of waiting time among different regions of a workspace.
  - ▶ Ensure a balanced exploration of the environment



Wu L., Puig D., and Garcia M. A. *Balanced Multi-Robot Exploration through a Global Optimization Strategy*. Journal of Physical Agents. 2010.

# Multi-Robot Exploration



- K-means is used to partition an unexplored space in regions.
- Each region is assigned to a robot solving an optimization problem.



Wu L., Puig D., and Garcia M. A. *Balanced Multi-Robot Exploration through a Global Optimization Strategy*. Journal of Physical Agents. 2010.

# Outdoor Robots Navigation


- Task definition
  - ▶ Autonomous navigation for outdoor, unstructured environments.
  - ▶ Recognize navigable terrain and avoid obstacles, based on their appearance.
- Problem
  - ▶ Do reliable segmentation of outdoor scenes in an efficient manner (online).



Blas M. R., Agrawal M., Sundaresan A., and Konolige K. *Fast Color/Texture Segmentation For Outdoor Robots*. IEEE International Conference on Intelligent Robots and Systems. 2008.



# Outdoor Robots Navigation



a)

b)


- Use compact texture/color descriptors (feature vectors) and fast unsupervised clustering algorithms.
- K-means is used to cluster neighborhood feature vectors in a small set of basis vectors (*textons*).
- Each pixel is classified as belonging to one cluster using Euclidean distance (b).




Blas M. R., Agrawal M., Sundaresan A., and Konolige K. *Fast Color/Texture Segmentation For Outdoor Robots*. IEEE International Conference on Intelligent Robots and Systems. 2008.



# Outdoor Robots Navigation



c)



d)

- Extract an histogram counting textons in a neighborhood of each pixel
- K-means to extract a set of histogram profiles ( $k = 8$ ).
- *Earth Movers Distance* is used to merge similar clusters (c).



Blas M. R., Agrawal M., Sundaresan A., and Konolige K. *Fast Color/Texture Segmentation For Outdoor Robots*. IEEE International Conference on Intelligent Robots and Systems. 2008.

# Summary and Next Lecture

- What we have learned
  - k-means algorithm
  - hierarchical clustering
- Reading: Duda Chap. 10.4, 10.6-9, Bishop Chap. 9.1, Michell Chap. 6.12
- Next Lecture
  - Maximum Likelihood Estimation
  - Gaussian Mixture Model