

# Machine Learning in Robotics

## Lecture 11: Introduction to Reinforcement Learning

**Prof. Dongheui Lee**

*Institute of Automatic Control Engineering  
Technische Universität München*

dhlee@tum.de


# Reinforcement Learning

Learning of a behavior without explicit information about correct actions

- Between supervised and unsupervised learning
- No training patterns, but **rewards**
- Inspired by principles of human and animal learning
- Mild assumptions on the process to be controlled
- A control strategy can be learned from scratch

# Architecture

The agent-environment interaction in reinforcement learning



# The Environment

- The environment contains the process to be controlled
- Markov Decision Process (MDP): The environment is modeled by an MDP which is tuple  $(S, A, \{P_{sa}\}, \gamma, R)$ 
  - ▶  $S$  is a set of **states**
  - ▶  $A$  is a set of **actions**
  - ▶  $P_{sa}$  are the **state transition probabilities**.
  - ▶  $\gamma \in [0, 1)$  is the **discount factor**.
  - ▶  $R : S \times A \mapsto \mathbb{R}$  is the **reward function** (Rewards can also be a function of state  $S$  only and in that case  $R : S \mapsto \mathbb{R}$ ).

# Markov Process, HMM, Markov Decision Process

# Task for the Agent

Find a behavior which maximizes the expected total reward

For how long should we consider?

## Finite Horizon

$$\max \left[ \sum_{t=0}^T r_t \right]$$

## Infinite Horizon

$$\max \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

$\gamma$  is a discount factor ( $0 \leq \gamma < 1$ )

# Reward function

The reward function controls which task should be solved

- Game (Checkers, chess)  
Reward only at end: +1 when winning, -1 when loosing
- Avoiding mistakes (pole balancing)  
Reward -1 at the end (when falling)
- Find a fast/short/cheap path to a goal  
Reward -1 at each step

# Simplifying assumptions

- Discrete time
- Finite number of actions  $a_i \in a_1, a_2, a_3, \dots, a_n$
- Finite number of states  $s_i \in s_1, s_2, s_3, \dots, s_m$
- Environment is a stationary markov decision process
- Reward  $r$  only depends on  $s$



# Policy and Value function

- Policy

Policy provides a mapping from states to action.

$$\pi(s) \mapsto a$$

- Value Function

Expected total future reward when starting from  $s$  and following policy  $\pi$

$$\begin{aligned} V^\pi(s) &= E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi] \\ &= R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s') \quad (\text{Bellman's equation}) \end{aligned}$$

# Optimal Policy

An optimal policy is the the one which maximizes the value function

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

$$\pi^*(s) = \arg \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

# Classical problem: Grid world

- 11 states. Each **state** is represented by a position in the grid world.
- The agent **acts** deterministically by moving to other position.  
 $A=\{N,S,E,W\}$
- reward:  $R(4,3) = 1$ ,  $R(4,2) = -1$ ,  $R(s) = -0.02$  for all other states
- transition probability: 0.8 for a planned state and 0.1 for the other adjacent two states.

# Value Iteration

For each state  $s$ , initialize  $V(s) := 0$ .  
Repeat until convergence  
{  
For every state, update  $V(s) := R(s) + \max_{a \in A} \gamma \sum_{s'} P_{sa}(s') V(s')$   
}

$V(s)$  can be updated in synchronous and asynchronous manner.

# Policy Iteration

Initialize  $\pi$  randomly.  
Repeat until convergence  
{  
(a) Let  $V := V^\pi$   
(b) For each state  $s$ , let  $\pi(s) := \arg \max_{a \in A} \sum_{s'} P_{sa}(s') V(s')$   
}

Step (a) can be calculated by solving linear equations (with equal number of equations and unknowns).

# Monte-Carlo Method

Start at some random state.

Follow  $\pi$ , store the rewards and  $s_t$ .

When the goal is reached, update  $V^\pi(s)$  estimation for all visited states with the future reward we actually received.

- Monte-Carlo method is suitable only for episodic tasks
- Learns incrementally from episode-by-episode but not step-by-step

# Temporal Difference Learning

There are two estimates of the value of a state:

- Before:  $V^\pi(s_t)$
- After:  $R_{t+1} + \gamma V^\pi(s_{t+1})$

# Temporal Difference Learning


Idea: The second estimate is better!

$$V^{\pi}(s_t) \leftarrow V^{\pi}(s_t) + \alpha(R_{t+1} + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t))$$

- Learns considerably faster than the Monte-Carlo method
- Step by step learning.



# Eligibility Trace



# Q-Learning

Whenever reward  $r$  or next state  $s'$  cannot be predicted, we cannot calculate  $\pi$  even with a good estimate for  $V$

$Q^\pi(s, a)$ , is the expected infinite-horizon discounted return for executing  $a$  in state  $s$  and thereafter following  $\pi$

$$Q^\pi(s, a) = E \left[ R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid s_t = s, a_t = a, \pi \right]$$

# Q-Learning

Whenever reward  $r$  or next state  $s'$  cannot be predicted, we cannot calculate  $\pi$  even with a good estimate for  $V$

$Q^\pi(s, a)$ , is the expected infinite-horizon discounted return for executing  $a$  in state  $s$  and thereafter following  $\pi$

$$Q^\pi(s, a) = E [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s_t = s, a_t = a, \pi]$$

$$\pi(s) = \arg \max_a Q(s, a)$$

$$V^*(s) = \max_a Q^*(s, a)$$

# Q-Learning

Whenever reward  $r$  or next state  $s'$  cannot be predicted, we cannot calculate  $\pi$  even with a good estimate for  $V$

$Q^\pi(s, a)$ , is the expected infinite-horizon discounted return for executing  $a$  in state  $s$  and thereafter following  $\pi$

$$Q^\pi(s, a) = E \left[ R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid s_t = s, a_t = a, \pi \right]$$

$$\pi(s) = \arg \max_a Q(s, a)$$

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

# Autonomous helicopter flight via RL




*Value* of each policy is calculated through Monte-Carlo method *PEGASUS* method uses the observation that almost all computer simulations sample  $s' \sim P_{sa}(\cdot)$  by first calling a random number generator to get one (or more) random numbers  $p$ , and then calculating  $s'$  as some deterministic function of the input  $s, a$  and the random  $p$  Since the helicopters model is stochastic, random number were fixed in advanced to evaluate different policies



HJ Kim, Michael I Jordan, Shankar Sastry, and Andrew Y Ng., *Autonomous helicopter flight via reinforcement learning*, In Advances in neural information processing systems, 2003.



# Learning Motor Primitives using Reinforcement Learning



*POWER* is an Expectation Maximization based RL algorithm which does not require learning rate as a parameter:

$$\theta' = \theta + \frac{E\{\sum_{t=1}^T \varepsilon_t Q^\pi(s_t, a_t, t)\}}{E\{\sum_{t=1}^T Q^\pi(s_t, a_t, t)\}} \quad \text{where } \varepsilon_t \text{ is exploration term}$$



Jens Kober and Jan Peters, *Learning motor primitives for robotics*, pp. 2112 - 2118, ICRA, 2009.



# Reading Material

- Mitchell, Chapter 13
- Russell and Norvig, Artificial Intelligence: A Modern Approach, Chapter 21
- Sutton and Barto, Reinforcement Learning: An Introduction